

ZIDANE VIRANI

647-667-1235 | zidanevirani54@gmail.com | [linkedin.com/in/zidanevirani](https://www.linkedin.com/in/zidanevirani) | [Portfolio](#)

Education

Western University

Bachelor of Engineering in Software Engineering

Apr 2027

GPA: 3.8/4.0

- Courses: Data Structures & Algorithms, Operating Systems, Neural Networks, NLP, Cloud Computing, Distributed Systems, Database Systems, Computer Networks, Software Architecture, Parallel Computing, Computer Vision, Statistical Learning
- Involvements: Outreach chair of Western Dental Association, Director at Western AI, Academic Mentor for Freshmen

Skills

- Python, C/C++, Go, Swift, JavaScript, SQL, Bash, Node.js, React, FastAPI, Django, Redis, Hugging Face, LangChain, CI/CD
- CUDA, AWS, Azure, PyTorch, Tensorflow, Keras, Docker, Kubernetes, GCP, PostgreSQL, Kafka, ONNX, Triton, OpenCV, FAISS

Experience

Founder & Lead Engineer, AI Content Engine ([Blog Link](#))

Jan 2023 – Dec 2024

- Sold a Python Retrieval Augmented Generation (RAG) pipeline for \$50,000 that scraped trending topics and auto-generated SEO-optimized articles with GPT-4, scaling my blog to 120K+ monthly readers and \$2,800/month in net profit
- Developed and deployed an end-to-end production-ready ML ranking model that increased article engagement by 412% by training on 10,000+ user-labeled examples using PyTorch, CUDA, Hugging Face Transformers, and multi-GPU EC2 infrastructure
- Reduced infrastructure costs by 57% by fine-tuning and deploying open-source LLM for SEO workflows and deploying it on EC2 Spot, building an inference pipeline that successfully processed all workloads previously done by GPT-4

Bloomberg : Incoming Machine Learning Engineer Intern

Jun 2026 - Aug 2026

Shopify: Applied Machine Learning Engineer Intern

Jan 2026 - Apr 2026

- Built an autocomplete feature for Shopify Sidekick using Python, PyTorch, Hugging Face Transformers, and a fine-tuned Qwen-0.6B LLM, reducing merchant query time by 68% and deploying to 5 Million+ active stores through Docker and Kubernetes
- Optimized Sidekick's ML inference pipeline using Python, PyTorch, ONNX, and Redis caching, cutting average response latency by 85%, saving 2,400+ GPU hours weekly and reducing infrastructure costs by 34% across production systems
- Engineered an automated data curation pipeline using PyTorch, SageMaker, and Spark to filter and rank 2M+ merchant queries through distributional analysis, boosting fine-tuned Qwen model eval performance by 3.8× on internal benchmarks

Node App: Machine Learning Developer Intern

Sep 2025 - Present

- Built an AI-powered search system (RAG) handling 5M+ documents using Python, Elasticsearch, and Docker, reducing response times by 92% through vector indexing and optimized caching techniques
- Improved search relevance by 91% by training a PyTorch ranking model on 500K+ click-through events and deploying it with FastAPI on AWS using Docker, Lambda, and DynamoDB, validating results through end-to-end A/B experiments
- Scaled the platform to handle 3× more traffic by using Docker for model containers, ONNX Runtime for faster inference, and Kubernetes autoscaling, reducing memory usage by 58% with sub-200ms latency

Western University: ML Engineer Intern

Jun 2025 – Sep 2025

- Led a team of 5 to build an LLM-powered fact-checking system using Python, Django, PyTorch, FAISS, Docker, and AWS EC2 Spot, and added a retrieval-and-verification pipeline that reached 92.4% precision across 1.2M academic papers
- Reduced hallucinations by 61% by integrating the system into GPT-3.5 pipelines as a verification layer to validate outputs
- Built a ML pipeline with Python, PyTorch, NLP techniques, and computer vision that combined text, audio, and visual cues to flag interruptions in political debates, deploying it to 70+ political science departments for live analysis and teaching

Vector Institute: Full-Stack ML Developer Intern

Apr 2025 - Jun 2025

- Built and deployed a senior-care voice assistant on AWS with Python, TensorFlow, Whisper and FastAPI, rolling it out across 18 retirement facilities (5,000+ users) and saving staff 1,200+ hours monthly through automated requests and daily care workflows
- Enabled multilingual support by fine-tuning OpenAI Whisper with PyTorch and LoRA on 50K+ caregiver–patient samples, boosting transcription accuracy by 34% and cutting clarification calls by 41% across all facilities

Dibbly Inc: AI Engineer Intern

Jan 2025 - Apr 2025

- Built a LLM-Powered chatbot (RAG) replacing a 3-day manual workflow for 5,000+ users, resolving 3,000+ weekly queries at 95% accuracy by deploying semantic search over 10,000+ documentation pages using LangChain and Pinecone vector database
- Optimized Mistral-7B (Open-source Large Language Model) fine-tuning in Python with CUDA, PyTorch, and AWS SageMaker, making training 28% faster by using more efficient fine-tuning techniques like LoRA and gradient checkpointing

Leadership

Aga Khan Development Network: Chair

Sep 2021 – Present

- Securing \$100,000 in funding through partnerships with Canadian healthcare leaders and spearhead the development of a women-empowerment platform in India using FastAPI, React, and AWS to build a digital ecosystem enabling fair-wage opportunities